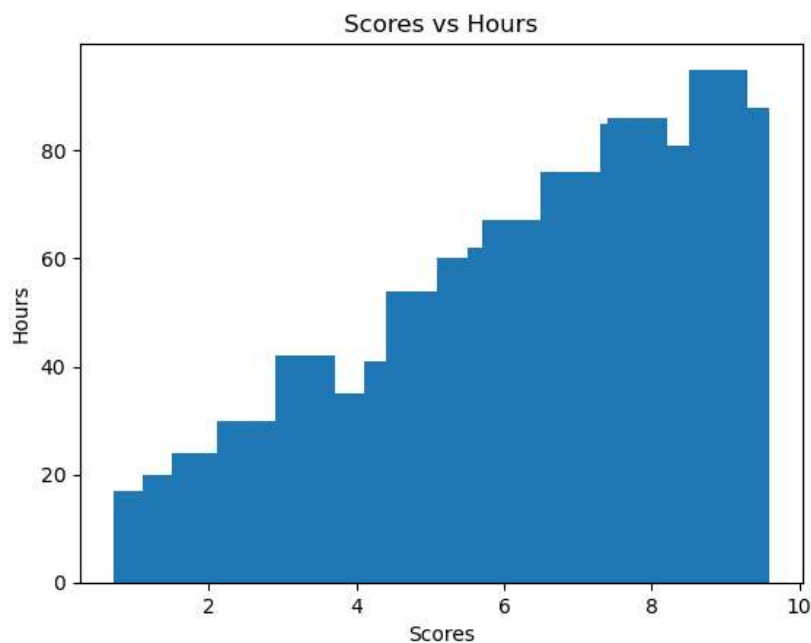



```
In [2]: 1 import pandas as pd
2 import numpy as np
3 from sklearn.model_selection import train_test_split
4 from sklearn.linear_model import LinearRegression
5 from sklearn.metrics import accuracy_score, mean_absolute_error, mean_squared_error
6 from sklearn.preprocessing import MinMaxScaler, StandardScaler
7 from sklearn.compose import ColumnTransformer
8 from sklearn.ensemble import RandomForestRegressor
9 import joblib
10 import matplotlib.pyplot as plt
11 student_df = pd.read_csv("student_scores.csv")
12 student_df
13 scores = student_df["Scores"]
14 hours = student_df["Hours"]
15 #bar plot
16 plt.bar(hours, scores)
17 #labels and title
18 plt.xlabel('Scores')
19 plt.ylabel('Hours')
20 plt.title('Scores vs Hours')
21 # show the plot
22 plt.show()
23
```



```
In [3]: 1 scores = student_df["Scores"]
2 hours = student_df["Hours"]
3 # Create the bar plot
4 plt.scatter(hours,scores);
5 plt.xlabel("Hours")
6 plt.ylabel("Scores")
7 plt.title("Hours vs Scores")
8
```

Out[3]: Text(0.5, 1.0, 'Hours vs Scores')



```
In [9]: 1 x = np.array(student_df["Hours"]).reshape(-1,1)
2 y = student_df["Scores"]
3 # Splitting the Data
4 x_train ,x_test ,y_train ,y_test = train_test_split(x,y,test_size=0.3, random_state=41)
5 regressor= LinearRegression()
6 regressor.fit(x_train ,y_train)
```

Out[9]:

LinearRegression

LinearRegression()

```
In [10]: 1 regressor.score(x_test,y_test)
```

Out[10]: 0.9621346134566173

```
In [21]: 1 joblib.dump(model, "StudentScore.pkl")
```

Out[21]: ['StudentScore.pkl']

```
In [22]: 1 model = joblib.load("StudentScore.pkl")
```

```
In [24]: 1 y_preds = model.predict(x_test)
          2 pd.DataFrame({"y_test": y_test, "y_pred": y_preds})
          3
```

Out[24]:

	y_test	y_pred
5	20	14.990287
19	69	74.884076
14	17	10.929691
10	85	77.929523
7	60	55.596245
8	81	84.020416
17	24	19.050883
11	62	59.656841

```
In [26]: 1 mean_square_error = mean_squared_error(y_test, y_preds)
          2 print(f"mean squared error is {mean_square_error}")
          3 maean_absolute_error = mean_absolute_error(y_test, y_preds)
          4 print(f"mean absolute error is {maean_absolute_error}")
          5
```

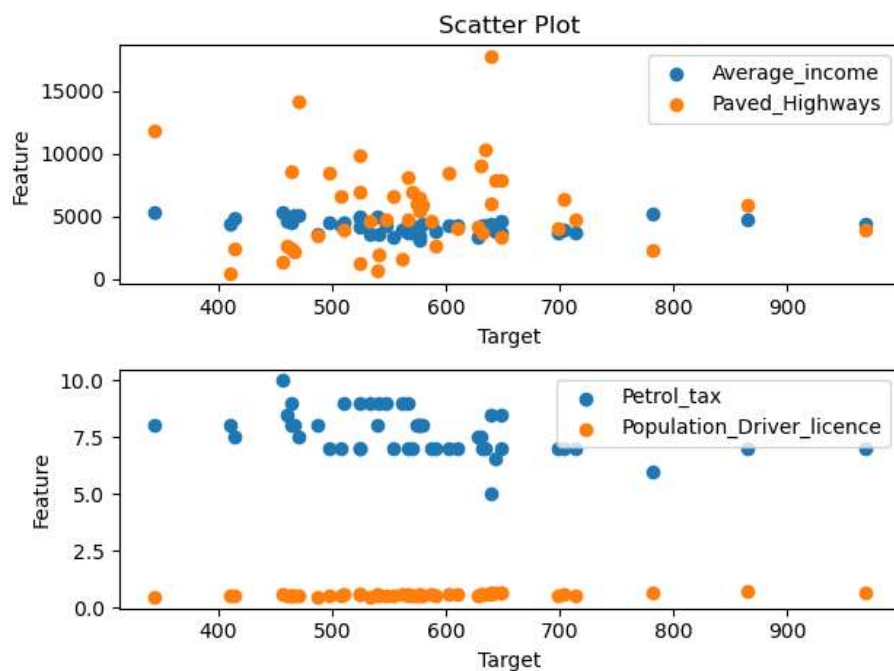
```
mean squared error is 25.63250010321114
mean absolute error is 4.843877805244228
```

Task#2

```

In [28]: 1 Petrol_df = pd.read_csv("petrol_consumption.csv")
2 Petrol_df.head()
3 target = Petrol_df['Petrol_Consumption']
4 Average_income = Petrol_df['Average_income']
5 Paved_Highways = Petrol_df['Paved_Highways']
6 Petrol_tax = Petrol_df['Petrol_tax']
7 Population_Driver_licence = Petrol_df['Population_Driver_licence(%)']
8 plt.subplot(2, 1, 1)
9 plt.scatter(target, Average_income, label='Average_income')
10 plt.scatter(target, Paved_Highways, label='Paved_Highways')
11 plt.xlabel('Target')
12 plt.ylabel('Feature')
13 plt.title('Scatter Plot')
14 plt.legend()
15 plt.subplot(2, 1, 2)
16 plt.scatter(target, Petrol_tax, label='Petrol_tax')
17 plt.scatter(target, Population_Driver_licence, label='Population_Driver_licence')
18 plt.xlabel('Target')
19 plt.ylabel('Feature')
20 plt.legend()
21 plt.tight_layout() # Adjusts the spacing between subplots
22 plt.show()
23

```



```

In [31]: 1 minmax =MinMaxScaler()
2 transform_columns=["Petrol_tax","Average_income","Paved_Highways","Population_Driver_licence(%)"]
3 transformer =ColumnTransformer([("minmax",minmax,transform_columns)])
4 data = transformer.fit_transform(Petrol_df)
5 data = pd.DataFrame(data,columns=["Petrol_tax","Average_income","Paved_Highways","Population_Driver_licence(%)"]
6 data.head()
7

```

```

Out[31]:

```

	Petrol_tax	Average_income	Paved_Highways	Population_Driver_licence(%)
0	0.8	0.222905	0.089044	0.271062
1	0.8	0.451514	0.047202	0.443223
2	0.8	0.351909	0.066567	0.472527
3	0.5	0.792892	0.110656	0.285714
4	0.6	0.586222	0.000000	0.340659

```

In [32]: 1 x = data
2 y = Petrol_df["Petrol_Consumption"]

```

```
In [33]: 1 x_train ,x_test ,y_train ,y_test = train_test_split(x,y,test_size=0.3 ,random_state=52)
2 model1 = LinearRegression()
3 model1.fit(x_train ,y_train)
```

```
Out[33]: ▾ LinearRegression
LinearRegression()
```

```
In [34]: 1 model1.score(x_test,y_test)
```

```
Out[34]: 0.7812556015921516
```

```
In [37]: 1 joblib.dump(model1, "PetrolConsumption_a.pkl")
2 mod1 = joblib.load("PetrolConsumption_a.pkl")
3 y_pred = mod1.predict(x_test)
4 pd.DataFrame({"y_test": y_test, "y_pred":y_pred})
```

```
Out[37]:
```

	y_test	y_pred
8	464	498.887986
3	414	488.882860
25	566	549.211935
44	782	675.788092
18	865	745.446171
26	577	597.792995
6	344	337.080999
9	498	547.265145
34	487	520.150864
17	714	594.451750
1	524	561.120649
30	571	570.485001
0	541	540.348735
38	648	714.693951
21	540	559.453224

```
In [38]: 1 mean_s_error = mean_squared_error(y_test, y_pred)
2 print(f"Mean Squared Error: {mean_s_error}")
3 mean_a_error = mean_absolute_error(y_test, y_pred)
4 print(f"Mean Absolute Error: {mean_a_error}")
```

```
Mean Squared Error: 3812.248209532195
Mean Absolute Error: 47.09566602370284
```

```
In [ ]: 1
```