

Usama Arif roll no 14

Predict the Stroke in the given data set and apply all available kernels in SVC model prepare the data set according to need (numeric)let us know the which kernel is best for such applicationWrite Story Telling

In [10]:

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import seaborn as sns
4 from sklearn.preprocessing import MinMaxScaler,OneHotEncoder,StandardScaler
5 from sklearn.compose import ColumnTransformer
6 from sklearn.model_selection import train_test_split
7 from sklearn.metrics import confusion_matrix,f1_score
8 from sklearn.svm import SVC
9 import numpy as np
```

In [11]:

```
1 data=pd.read_csv('healthcare.csv')
2 data=data.drop('id',axis=1)
3 data.head()
```

Out[11]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	avg_
0	Male	67.0	0	1	Yes	Private	Urban	
1	Female	61.0	0	0	Yes	Self-employed	Rural	
2	Male	80.0	0	1	Yes	Private	Rural	
3	Female	49.0	0	0	Yes	Private	Urban	
4	Female	79.0	1	0	Yes	Self-employed	Rural	

In [12]:

```
1 data.isna().sum()
```

Out[12]:

```
gender          0
age             0
hypertension     0
heart_disease    0
ever_married     0
work_type        0
Residence_type   0
avg_glucose_level 0
bmi             201
smoking_status   0
stroke           0
dtype: int64
```


In [14]:

```
1 data=data.dropna()
2 data
```

Out[14]:

	gender	age	hypertension	heart_disease	ever_married	work_type	Residence_type	a
0	Male	67.0	0	1	Yes	Private	Urban	
2	Male	80.0	0	1	Yes	Private	Rural	
3	Female	49.0	0	0	Yes	Private	Urban	
4	Female	79.0	1	0	Yes	Self-employed	Rural	
5	Male	81.0	0	0	Yes	Private	Urban	
...	
5104	Female	13.0	0	0	No	children	Rural	
5106	Female	81.0	0	0	Yes	Self-employed	Urban	
5107	Female	35.0	0	0	Yes	Self-employed	Rural	
5108	Male	51.0	0	0	Yes	Private	Rural	
5109	Female	44.0	0	0	Yes	Govt_job	Urban	

4909 rows × 11 columns



Data separation

In [15]:

```
1 x=data.drop('stroke',axis=1)
2 y=data.stroke
```

In [18]:

```
1 one_hot=OneHotEncoder()
2 scaler=StandardScaler()
3 features=['gender','ever_married','work_type','Residence_type','smoking_status']
4 trans=['age','avg_glucose_level','bmi']
5 transformer=ColumnTransformer([('features',one_hot,features),('trans',scaler,trans)]
6                               remainder='passthrough')
7 transformed_x=transformer.fit_transform(x)
8 transformed_x
```

Out[18]:

```
array([[ 0.          ,  1.          ,  0.          , ...,  0.98134488,
         0.          ,  1.          ],
       [ 0.          ,  1.          ,  0.          , ...,  0.45926914,
         0.          ,  1.          ],
       [ 1.          ,  0.          ,  0.          , ...,  0.70120668,
         0.          ,  0.          ],
       ...,
       [ 1.          ,  0.          ,  0.          , ...,  0.21733161,
         0.          ,  0.          ],
       [ 0.          ,  1.          ,  0.          , ..., -0.41934612,
         0.          ,  0.          ],
       [ 1.          ,  0.          ,  0.          , ..., -0.34294479,
         0.          ,  0.          ]])
```

training the data

we have to try different kernels for svc {'linear', 'poly', 'rbf', 'sigmoid', 'precomputed'}

In [25]:

```
1 svc_rbf=SVC(kernel='rbf')
2 np.random.seed(4)
3 x_train,x_test,y_train,y_test=train_test_split(transformed_x,y,test_size=0.2,random_
4 svc_rbf.fit(x_train,y_train);
```

Evaluating the model

In [26]:

```
1 svc_rbf.score(x_test,y_test)
```

Out[26]:

0.9541751527494908

In [28]:

```
1 y_preds=svc_rbf.predict(x_test)
2 f1_score(y_test,y_preds)
```

Out[28]:

0.0

In [29]:

```
1 svc_lin=SVC(kernel='linear')
```

In [30]:

```
1 svc_lin.fit(x_train,y_train)
```

Out[30]:

```
▼      SVC
SVC(kernel='linear')
```

In [31]:

```
1 svc_lin.score(x_test,y_test)
```

Out[31]:

0.9541751527494908

```
1 y_Pred=svc_lin.predict(x_test)
2 y_Pred
3
```

[illegible]

In [35]:

```
1 f1_score(y_test,y_Pred)
```

Out[35]:

0.0

In [38]:

```
1 svc_poly=SVC(kernel='poly')  
2 svc_poly.fit(x_train,y_train);
```

In [39]:

```
1 svc_poly.score(x_test,y_test)
```

Out[39]:

0.9541751527494908

```
1 y__pred=svc_poly.predict(x_test)
2 y__pred
3
```

[illegible]

In [47]:

```
1 f1_score(y__pred,y_test)
```

Out[47]:

0.0

In [48]:

```
1 svc_sigmoid=SVC(kernel='sigmoid')
2 svc_sigmoid.fit(x_train,y_train);
```

In [49]:

```
1 svc_sigmoid.score(x_test,y_test)
```

Out[49]:

0.9195519348268839


```
1 y_preds=svc_sigmoid.predict(x_test)
2 y_preds
```

[illegible]

```
1 f1_score(y__preds,y_test)
```

0.15053763440860213

In [57]:

```
1 svc_pre=SVC(kernel='precomputed')
2 svc_pre.fit(x_train,y_train)
```

```
-----
-
ValueError                                Traceback (most recent call las
t)
Cell In[57], line 2
      1 svc_pre=SVC(kernel='precomputed')
----> 2 svc_pre.fit(x_train,y_train)

File ~\anaconda3\lib\site-packages\sklearn\svm\_base.py:217, in BaseLibSV
M.fit(self, X, y, sample_weight)
    211     raise ValueError(
    212         "X and y have incompatible shapes.\n"
    213         + "X has %s samples, but y has %s." % (n_samples, y.shape
[0]))
    214 )
    216 if self.kernel == "precomputed" and n_samples != X.shape[1]:
--> 217     raise ValueError(
    218         "Precomputed matrix must be a square matrix."
    219         " Input is a {}x{} matrix.".format(X.shape[0], X.shape[1])
    220     )
    222 if sample_weight.shape[0] > 0 and sample_weight.shape[0] != n_samp
les:
    223     raise ValueError(
    224         "sample_weight and X have incompatible shapes: "
    225         "%r vs %r\n"
    (... )
    228         % (sample_weight.shape, X.shape)
    229     )
```

ValueError: Precomputed matrix must be a square matrix. Input is a 3927x21 matrix.

given data is not in sqaure form so precomputed kernel is not aplicable

In []:

```
1
```